

Seminar 3

Exploring Disease Classification Methods for Volumetric (3D) Medical Image (CTs/MRIs)



AI VIET NAM
[@aivietnam.edu.vn](http://aivietnam.edu.vn)



Presenter: Vũ Mai Anh
Upcoming PhD Student (Fall 2024)
University of Houston

Outline

Volumetric data

3D CNN

3D Unet

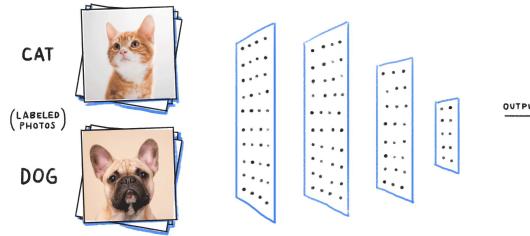
Exploring RSNA 2023 with 3D Unet and 2.5D CNN + RNN

The screenshot shows a competition card for the RSNA 2023 Abdominal Trauma Detection challenge. At the top left is the RSNA logo. To its right, the text reads "RADIOLOGICAL SOCIETY OF NORTH AMERICA · FEATURED CODE COMPETITION · 9 MONTHS AGO". On the far right are buttons for "Late Submission" and three vertical dots. The main title "RSNA 2023 Abdominal Trauma Detection" is centered in large white font. Below it, the subtitle "Detect and classify traumatic abdominal injuries" is written in smaller white font. To the right of the text is a small thumbnail image showing a grayscale axial CT scan of a human abdomen.

Introduction

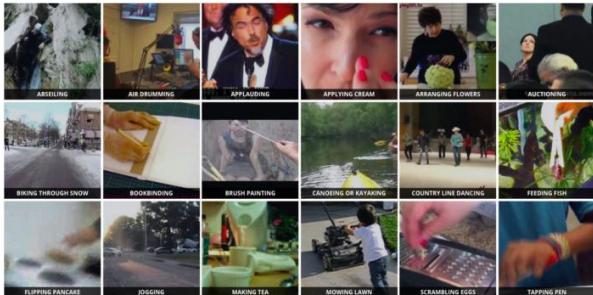
□ Volumetric Data

Image

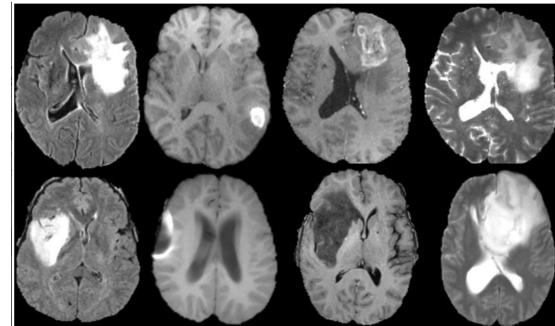


Sources [OleksandrKosovan](#)

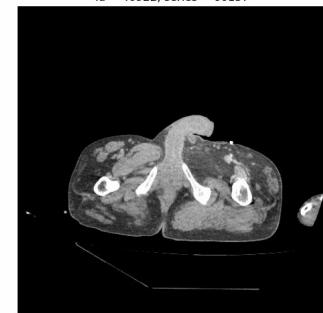
Series of Images



Kinetics Human Action Video Dataset
(Kay et al.)



The Multimodal Brain Tumor Image Segmentation Benchmark
(BRATS 2017, Bjoern H. Menze et al.)

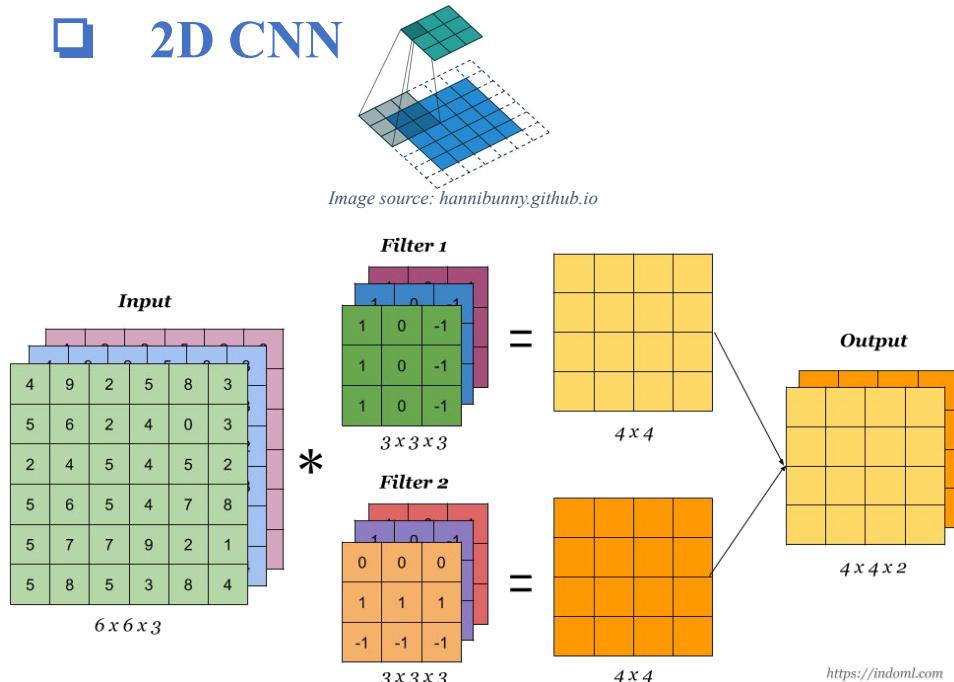


Frame 1

RSNA Abdominal Trauma Detection AI
Challenge (2023)

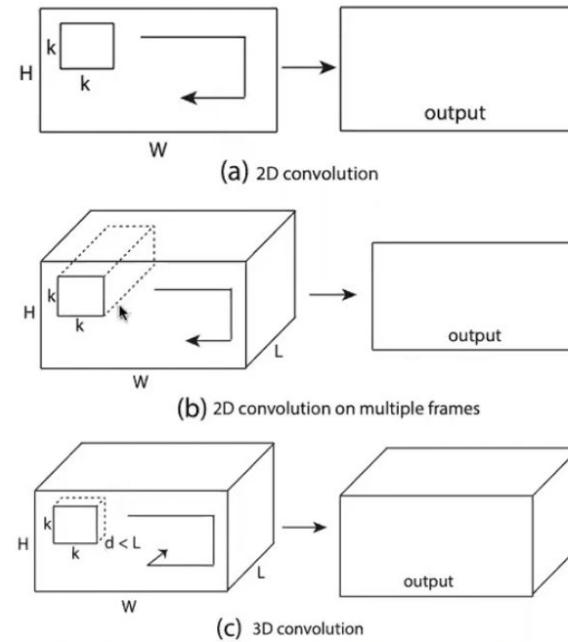
2D CNN to 3D CNN

2D CNN



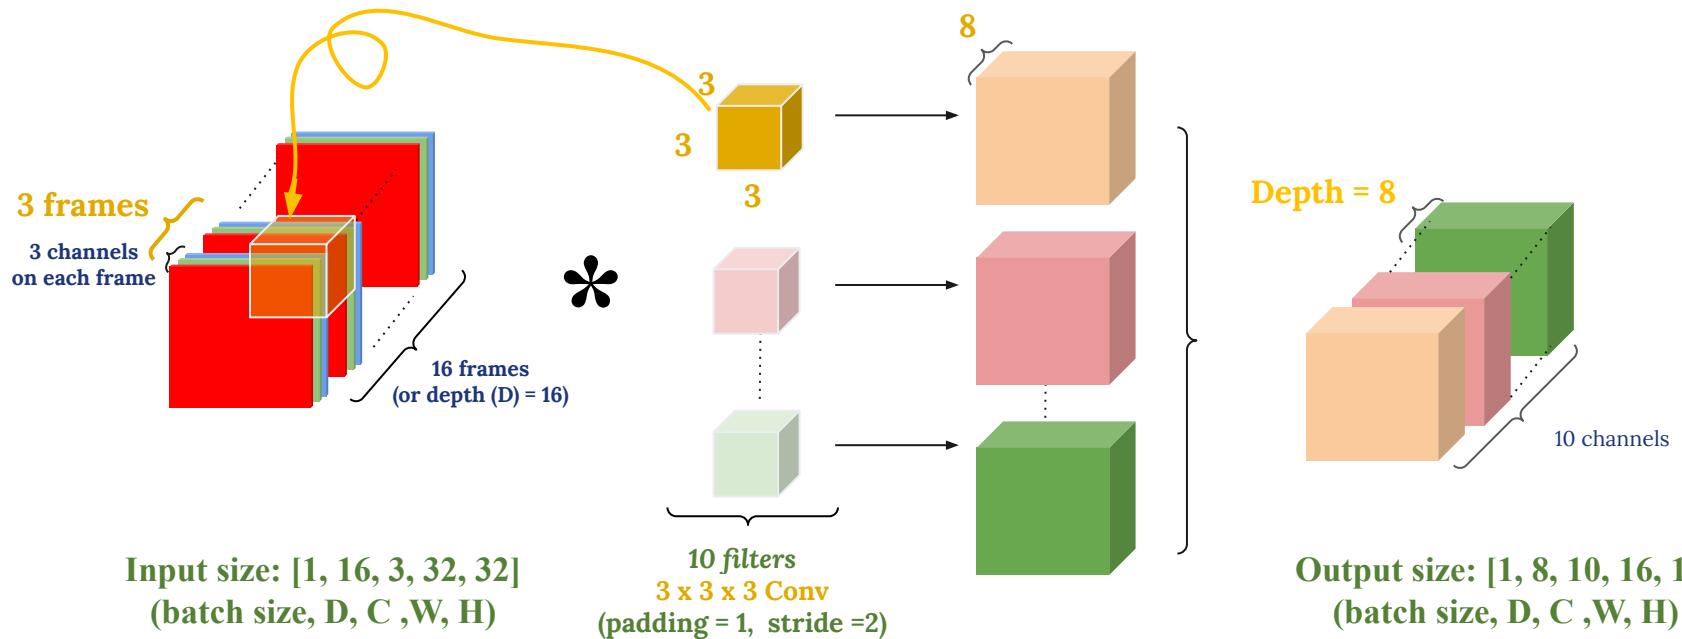
Source: <https://www.coursera.org/learn/convolutional-neural-networks>

2D CNN to 3D CNN



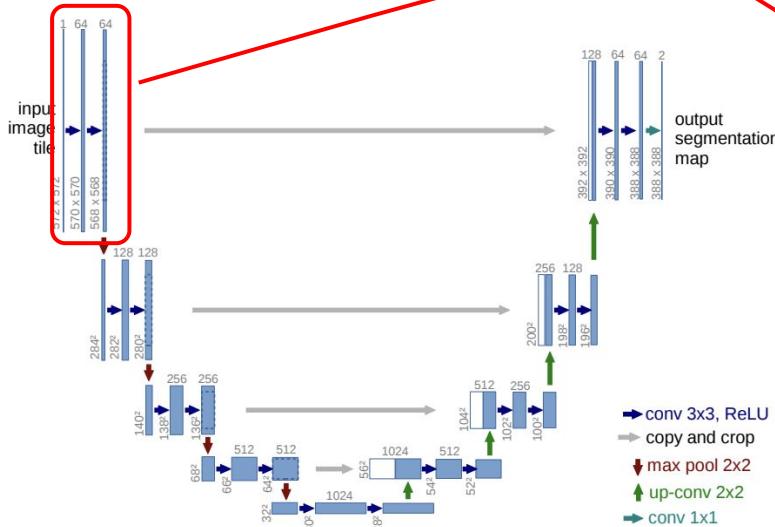
2D CNN to 3D CNN

3D CNN



2D Unet to 3D Unet

Review 2D Unet



```
class DoubleConv(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(DoubleConv, self).__init__()
        self.conv = nn.Sequential(
            nn.Conv2d(in_channels, out_channels, 3, 1, 1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True),

            nn.Conv2d(out_channels, out_channels, 3, 1, 1, bias=False),
            nn.BatchNorm2d(out_channels),
            nn.ReLU(inplace=True)
        )
    def forward(self, x):
        return self.conv(x)
```

```
class UNET2D(nn.Module):
    def __init__(self, in_channels=3, out_channels=1, features = [64, 128, 256, 512]):
        super(UNET2D, self).__init__()
        self.downsamplings = nn.ModuleList()
        self.upsamplings = nn.ModuleList()
        self.pool = nn.MaxPool2d(2, 2)

        #down sampling
        for feature in features:
            self.downsamplings.append(DoubleConv(in_channels, feature))
            in_channels = feature

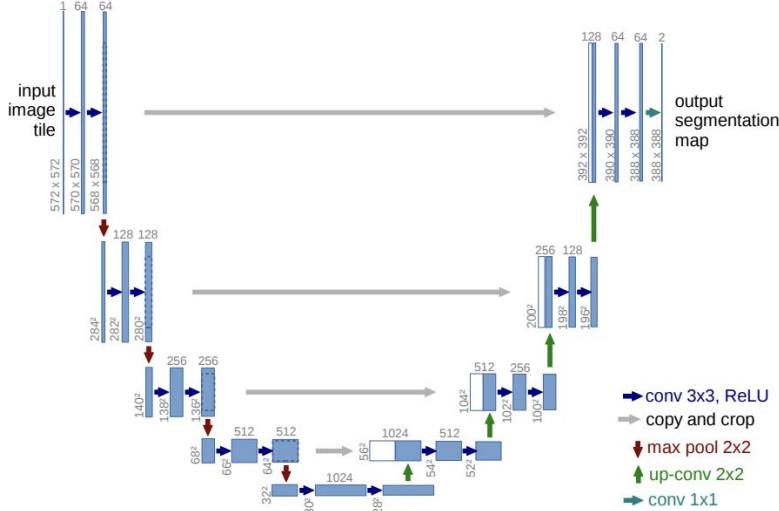
        #up sampling
        for feature in reversed(features):
            self.upsamplings.append(nn.Sequential(
                nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True),
                DoubleConv(feature * 2, feature)
            ))
            self.upsamplings.append(DoubleConv(feature*2, feature))

        self.bottleneck = DoubleConv(features[-1], features[-1]*2)
        self.final_conv = nn.Conv2d(features[0], out_channels, kernel_size=1)
```

U-Net: Convolutional Networks for Biomedical Image Segmentation

2D Unet to 3D Unet

Review 2D Unet



U-Net: Convolutional Networks for Biomedical Image Segmentation

Olaf Ronneberger, Philipp Fischer, and Thomas Brox

Computer Science Department at BIOSS Centre for Biological Signalling Studies,
University of Freiburg, Germany
ronneber@informatik.uni-freiburg.de
WWW home page: <http://lmb.informatik.uni-freiburg.de/>

```
class UNET2D(nn.Module):
    def __init__(self, in_channels=3, out_channels=1, features = [64, 128, 256, 512]):
        super(UNET2D, self).__init__()
        self.downsamplings = nn.ModuleList()
        self.upsamplings = nn.ModuleList()
        self.pool = nn.MaxPool2d(2, 2)

        #down sampling
        for feature in features:
            self.downsamplings.append(DoubleConv(in_channels, feature))
            in_channels = feature

        #up sampling
        for feature in reversed(features):
            self.upsamplings.append(nn.Sequential(
                nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True),
                DoubleConv(feature * 2, feature)
            ))
            self.upsamplings.append(DoubleConv(feature*2, feature))

        self.bottleneck = DoubleConv(features[-1], features[-1]*2)
        self.final_conv = nn.Conv2d(features[0], out_channels, kernel_size=1)

    def forward(self, x):
        skip_connection = []
        for down_sampling in self.downsamplings:
            x = down_sampling(x)
            skip_connection.append(x)
            x = self.pool(x)

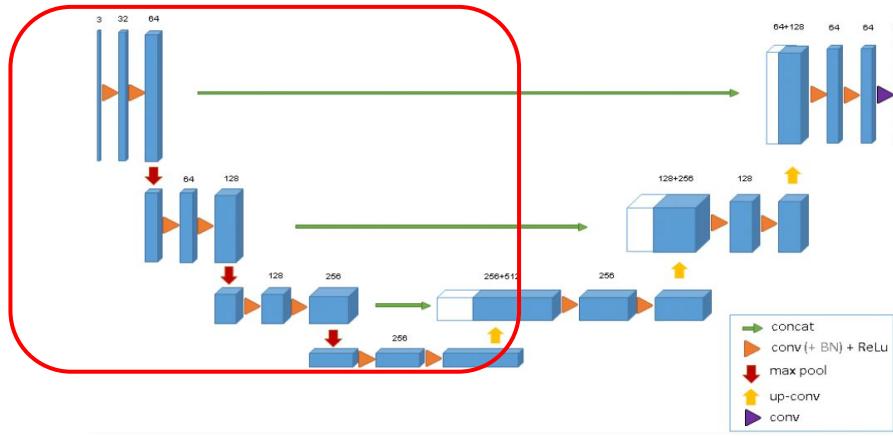
        x = self.bottleneck(x)
        skip_connection = skip_connection[::-1] #reverse

        for idx in range(0, len(self.upsamplings), 2):
            x = self.upsamplings[idx](x)
            skip_connection_feature = skip_connection[idx//2]
            concatenated_features = torch.cat([x, skip_connection_feature], dim=1)
            x = self.upsamplings[idx+1](concatenated_features)

        return self.final_conv(x)
```

2D Unet to 3D Unet

3D Unet



```
● ● ●

class DownSamplingDouble3DCNN(nn.Module):
    def __init__(self, in_channels=None, out_channels=None, bottleneck=False) → None:
        super(DownSamplingDouble3DCNN, self).__init__()
        self.conv1 = nn.Conv3d(
            in_channels=in_channels, out_channels=out_channels//2, kernel_size=(3,3,3), padding=1)
        self.bn1 = nn.BatchNorm3d(num_features=out_channels//2)
        self.conv2 = nn.Conv3d(
            in_channels=out_channels//2, out_channels=out_channels, kernel_size=(3,3,3), padding=1)
        self.bn2 = nn.BatchNorm3d(num_features=out_channels)
        self.relu = nn.ReLU()
        self.bottleneck = bottleneck
        if not bottleneck:
            self.pooling = nn.MaxPool3d(kernel_size=(2,2,2), stride=2)

    def forward(self, x):
        residual = self.relu(self.bn1(self.conv1(x)))
        residual = self.relu(self.bn2(self.conv2(residual)))
        out = residual
        if not self.bottleneck:
            out = self.pooling(residual)
        return out, residual
```

3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation

Özgün Çiçek^{1,2}, Ahmed Abdulkadir^{1,4}, Soeren S. Lienkamp^{2,3}, Thomas Brox^{1,2}, and Olaf Ronneberger^{1,3,5}

¹ Computer Science Department, University of Freiburg, Germany

² BROSS Centre for Biological Signalling Studies, Freiburg, Germany

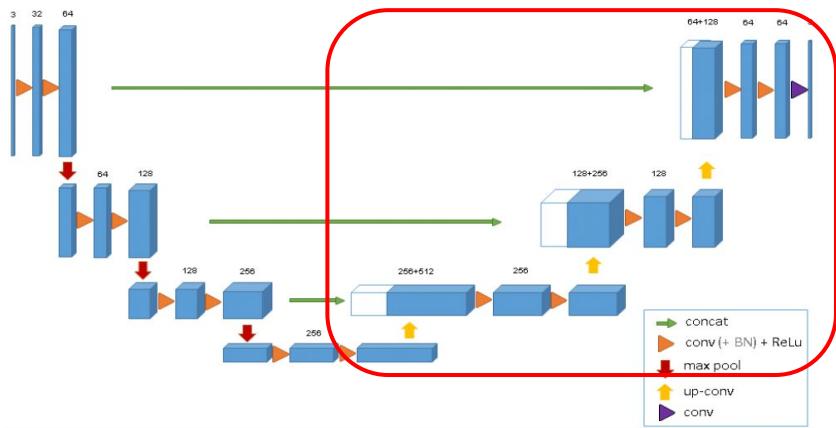
³ University Hospital Freiburg, Renal Division, Faculty of Medicine, University of Freiburg, Germany

⁴ Department of Psychiatry and Psychotherapy, University Medical Center Freiburg, Germany

⁵ Google DeepMind, London, UK
cicek@cs.uni-freiburg.de

2D Unet to 3D Unet

3D Unet



3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation

Özgür Çiçek^{1,2}, Ahmed Abdulkadir^{1,4}, Soeren S. Lienkamp^{2,3}, Thomas Brox^{1,2}, and Olaf Ronneberger^{1,2,5}

¹ Computer Science Department, University of Freiburg, Germany

² BIOSS Centre for Biological Signalling Studies, Freiburg, Germany
³ University Hospital Freiburg, Renal Division, Faculty of Medicine, University of Freiburg, Germany

⁴ Department of Psychiatry and Psychotherapy, University Medical Center Freiburg, Germany

⁵ Google DeepMind, London, UK
cicek@cs.uni-freiburg.de

```
class UpsamplingDouble3DCNN(nn.Module):
    def __init__(self, in_channels=None, res_channels=0, out_channels=None, last_layer=False,
                 num_classes=None):
        super(UpsamplingDouble3DCNN, self).__init__()
        self.last_layer = last_layer
        self.num_classes = num_classes

        self.upconv1 = nn.ConvTranspose3d(
            in_channels=in_channels, out_channels=in_channels, kernel_size=(2, 2, 2), stride=2)
        self.relu = nn.ReLU()
        self.bn1 = nn.BatchNorm3d(num_features=in_channels)
        self.bn2 = nn.BatchNorm3d(num_features=in_channels // 2)
        self.conv1 = nn.Conv3d(
            in_channels=in_channels + res_channels,
            out_channels=in_channels,
            kernel_size=(3, 3, 3),
            padding=1)
        self.conv2 = nn.Conv3d(
            in_channels=in_channels, out_channels=in_channels // 2, kernel_size=(3, 3, 3), padding=1)
        if self.last_layer:
            self.conv3 = nn.Conv3d(
                in_channels=in_channels // 2,
                out_channels=num_classes,
                kernel_size=(1, 1, 1))

    def forward(self, x, residual=None):
        out = self.upconv1(x)

        if residual is not None:
            if out.shape[2:] != residual.shape[2:]:
                residual = F.interpolate(
                    residual,
                    size=out.shape[2:],
                    mode='trilinear',
                    align_corners=False)

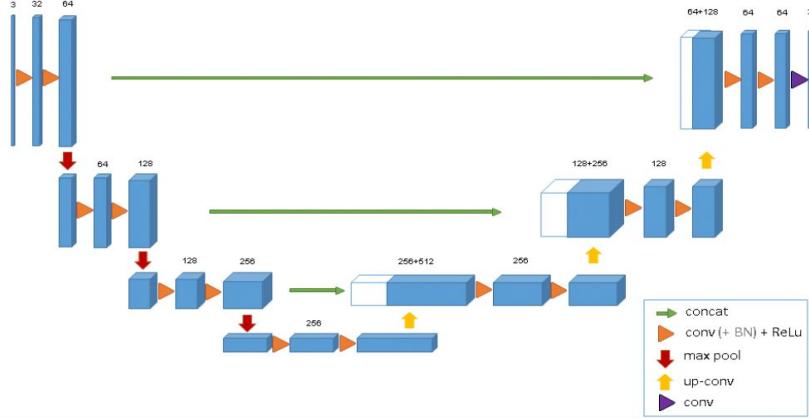
            out = torch.cat([out, residual], dim=1)

        out = self.relu(self.bn1(self.conv1(out)))
        out = self.relu(self.bn2(self.conv2(out)))

        if self.last_layer:
            out = self.conv3(out)
        return out
```

2D Unet to 3D Unet

3D Unet



```
class Unet3D(nn.Module):
    def __init__(self, in_channels=3, num_classes=3, num_features=[64, 128, 256, 512]):
        super(Unet3D, self).__init__()
        self.downsamplings = nn.ModuleList()
        self.upsamplings = nn.ModuleList()

    # Down sampling path
    for i in range(len(num_features) - 1):
        self.downsamplings.append(
            DownSampingDouble3DCNN(
                in_channels=in_channels if i == 0 else num_features[i-1],
                out_channels=num_features[i]))
    
    # Bottleneck
    self.bottleneck = DownSampingDouble3DCNN(
        in_channels=num_features[-2], out_channels=num_features[-1], bottleneck=True)

    # Up sampling path
    for i in range(len(num_features) - 1, 0, -1):
        self.upsamplings.append(
            UpsamplingDouble3DCNN(in_channels=num_features[i], res_channels=num_features[i-1],
                                  out_channels=num_features[i-1]))
    
    self.last_layer = UpsamplingDouble3DCNN(in_channels=num_features[0], res_channels=0,
                                             last_layer=True, num_classes=num_classes)

    def forward(self, x):
        residuals = []

        # Down sampling path
        for downsampling_layer in self.downsamplings:
            x, residual = downsampling_layer(x)
            residuals.append(residual)

        x, _ = self.bottleneck(x)

        # Up sampling path
        for idx, upsampling_layer in enumerate(self.upsamplings):
            x = upsampling_layer(x, residual=residuals[-idx-1])

        x = self.last_layer(x)
        return x
```

RSNA - RADIOLOGICAL SOCIETY OF NORTH AMERICA

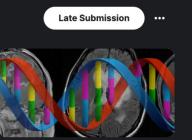
2021

RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 3 YEARS AGO

RSNA-MICCAI Brain Tumor Radiogenomic Classification

Predict the status of a genetic biomarker important for brain cancer treatment

Late Submission ...



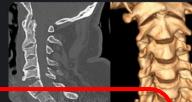
2022

RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 2 YEARS AGO

RSNA 2022 Cervical Spine Fracture Detection

Identify cervical fractures from scans

Late Submission ...



2023

RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 8 MONTHS AGO

RSNA 2023 Abdominal Trauma Detection

Detect and classify traumatic abdominal injuries

Late Submission ...



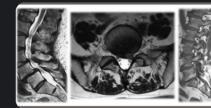
2024

RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 2 MONTHS TO GO

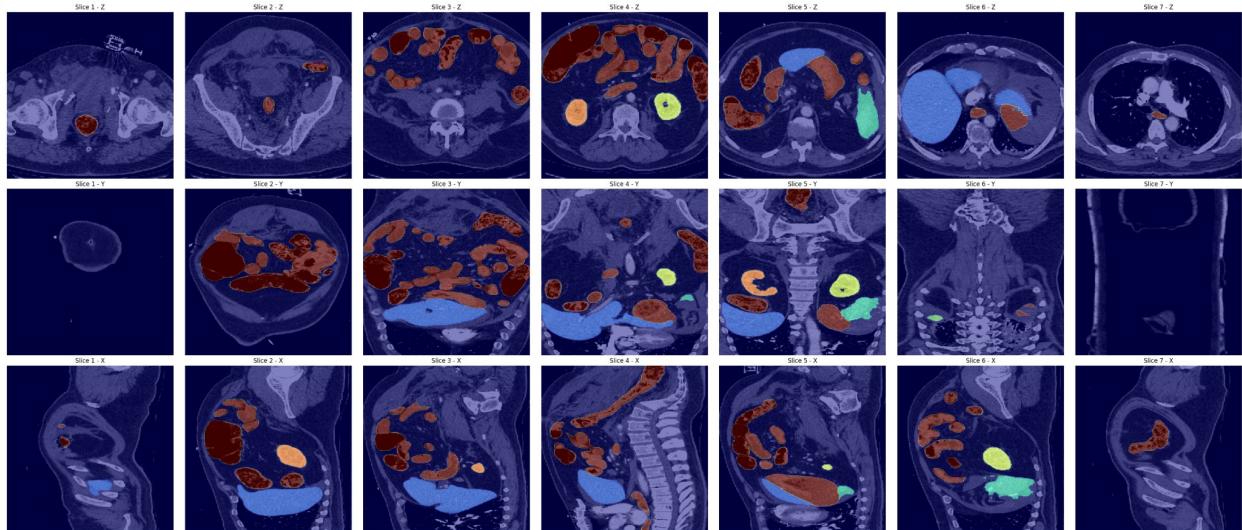
RSNA 2024 Lumbar Spine Degenerative Classification

Classify lumbar spine degenerative conditions

Submit Prediction ...



Volumetric data - Sample of 1 study (5 organs)



RNSA 2023

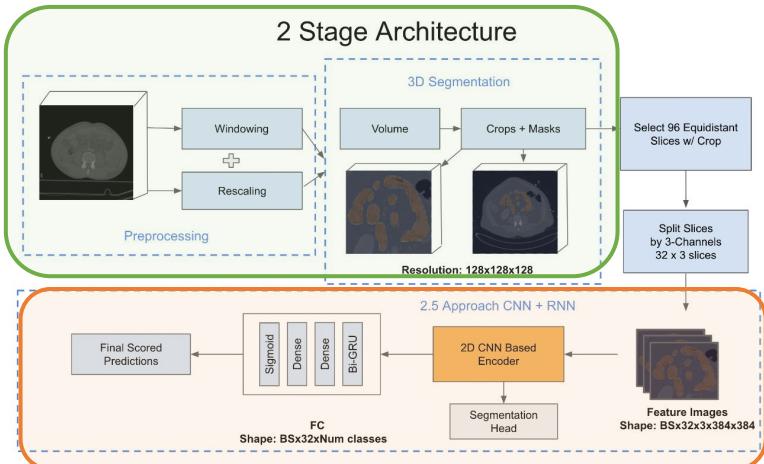
RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 9 MONTHS AGO

RSNA 2023 Abdominal Trauma Detection

Detect and classify traumatic abdominal injuries

Late Submission ...

Stage 1 - 3D segmentation

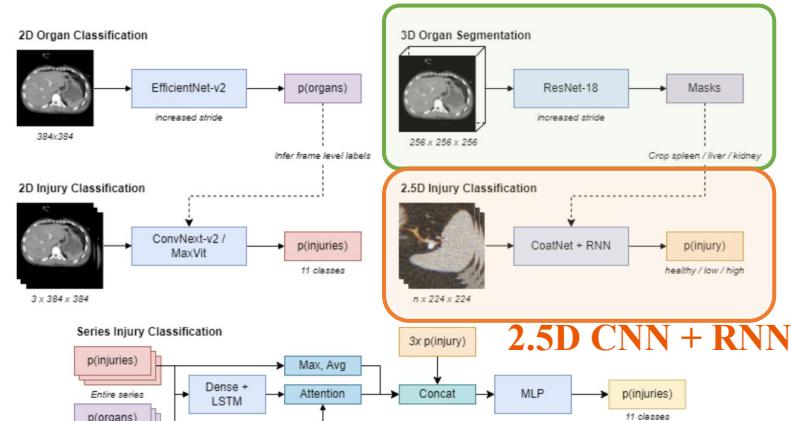


Stage 2 - 2.5D CNN + RNN

Source: 1st place solution:

<https://github.com/Nischaydnk/RSNA-2023-1st-place-solution>

3D segmentation



Source: 2nd place solution:
https://github.com/TheoViel/kaggle_rsna_abdominal_trauma

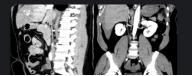
RNSA 2023

RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 9 MONTHS AGO

Late Submission ...

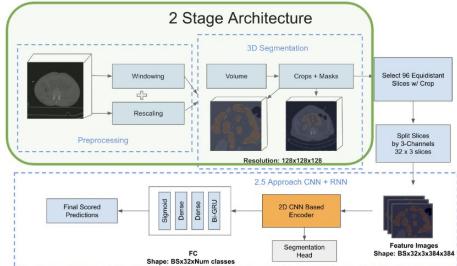
RSNA 2023 Abdominal Trauma Detection

Detect and classify traumatic abdominal injuries

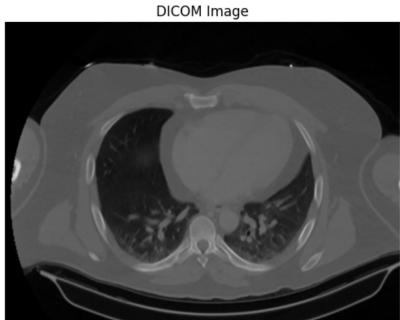


Preprocessing

Stage 1 - 3D segmentation



Stage 2 - 2.5D CNN + RNN



Windowing

```
[14]: import pandas as pd

# Specify the file path
file_path = '/kaggle/input/rsna-2023-abdominal-trauma-detection/train_dicom_tags.parquet'

# Display the first few rows of the DataFrame
print(dicom_tags_df.shape)
dicom_tags_df.head()

1500653, 35
signr ImageOrientationPatient ImagePositionPatient InstanceNumber ... WindowCenter WindowWidth FileMetaInformationVersion ImplementationClassUID ...
15 [1.0000, 0.0000, -235.8665, 1459.20] 532 ... 50.0 400.0 1.2.3.123456.4.5.1234.1.12.0
15 [1.0000, 0.0000, 1.0000, ... -235.8665, 1459.20] 513 ... 50.0 400.0 1.2.3.123456.4.5.1234.1.12.0
15 [1.0000, 0.0000, 1.0000, ... -235.8665, 1459.20] 718 ... 50.0 400.0 1.2.3.123456.4.5.1234.1.12.0
15 [1.0000, 0.0000, -235.8665, 1310.40] 456 ... 50.0 400.0 1.2.3.123456.4.5.1234.1.12.0
15 [1.0000, 0.0000, 1.0000, ... -235.8665, 1520.00] 161 ... 50.0 400.0 1.2.3.123456.4.5.1234.1.12.0
```

def get_windowed_image(dcm, WL=50, WW=400):
 resI, resS = dcm.RescaleIntercept, dcm.RescaleSlope

 img = dcm.to_numpy_image()

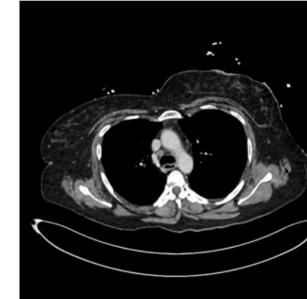
 img = get_standardized_pixel_array(dcm)

 img = resS * img + resI

 upper, lower = WL+WW//2, WL-WW//2
 X = np.clip(img.copy(), lower, upper)
 X = X - np.min(X)
 X = X / np.max(X)
 X = (X*255.0).astype('uint8')

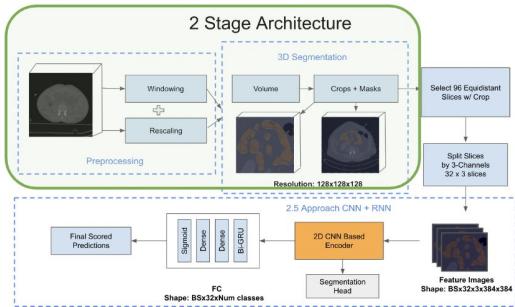
 return X

Image



Preprocessing

Stage 1 - 3D segmentation



Stage 2 - 2.5D CNN + RNN

Rescaling

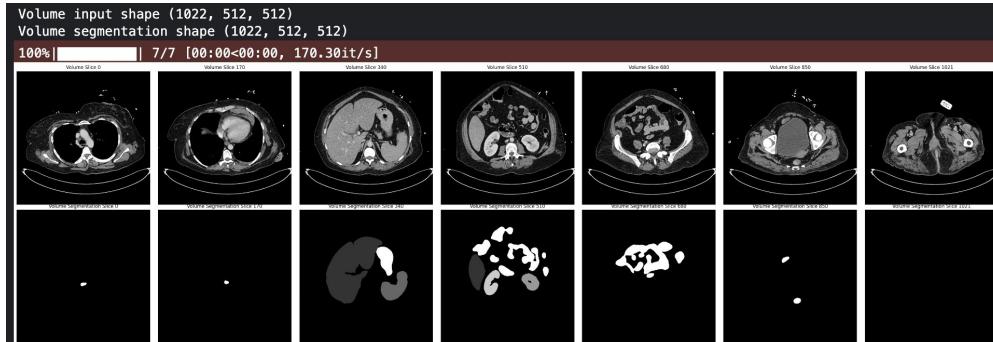
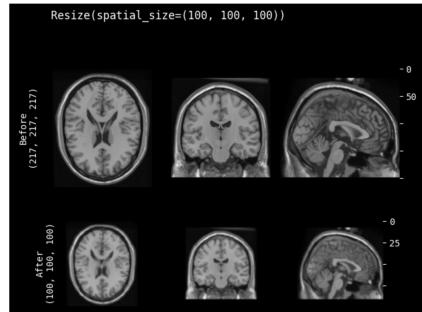


Figure: plot volume size and volume segment size



Source: monai

```
from monai.transforms import Resize
image_sizes = [128, 128, 128]
R = Resize(image_sizes)
```

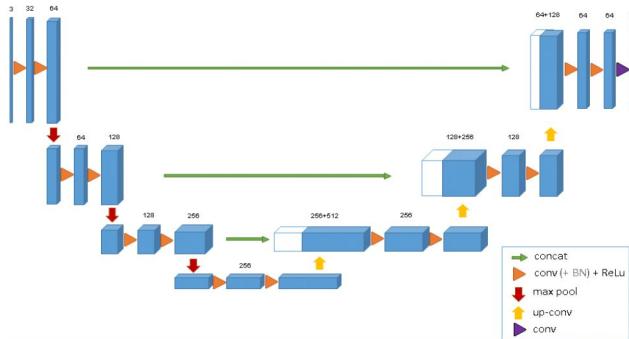
RNSA 2023

Late Submission ...

RSNA 2023 Abdominal Trauma Detection
Detect and classify traumatic abdominal injuries

3D segmentation

Stage 1 - 3D segmentation



```
class TimmSegModel(nn.Module):
    def __init__(self, backbone, segtype='unet', pretrained=False):
        super(TimmSegModel, self).__init__()

        self.encoder = timm.create_model(
            backbone,
            in_chans=3,
            features_only=True,
            drop_rate=drop_rate,
            drop_path_rate=drop_path_rate,
            pretrained=pretrained
        )
        g = self.encoder(torch.rand(1, 3, 64, 64))
        encoder_channels = [1] + [_.shape[1] for _ in g]
        decoder_channels = [256, 128, 64, 32, 16]
        if segtype == 'unet':
            self.decoder = smp.decoders.unet.decoder.UnetDecoder(
                #self.decoder = smp.unet.decoder.UnetDecoder(
                encoder_channels=encoder_channels[:n_blocks+1],
                decoder_channels=decoder_channels[:n_blocks],
                n_blocks=n_blocks,
            )

            self.segmentation_head = nn.Conv2d(
                decoder_channels[n_blocks-1],
                out_dim, kernel_size=(3, 3),
                stride=(1, 1), padding=(1, 1))

        def forward(self,x):
            global_features = [0] + self.encoder(x)[:n_blocks]
            seg_features = self.decoder(*global_features)
            seg_features = self.segmentation_head(seg_features)
            return seg_features
```

```
backbone = 'resnet18d'
m = TimmSegModel(backbone)
m = convert_3d(m)
print(m(torch.rand(1, 3, 128,128,128)).shape)
```

```
def convert_3d(module):

    module_output = module
    if isinstance(module, torch.nn.BatchNorm2d):
        ...

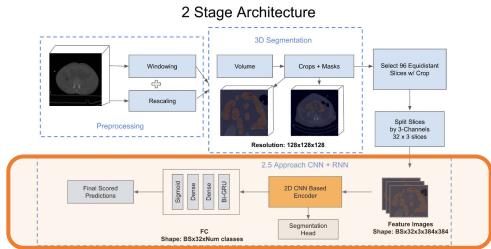
    elif isinstance(module, Conv2dSame):
        module_output = Conv3dSame(
            in_channels=module.in_channels,
            out_channels=module.out_channels,
            kernel_size=module.kernel_size[0],
            stride=module.stride[0],
            padding=module.padding[0],
            dilation=module.dilation[0],
            groups=module.groups,
            bias=module.bias is not None,
        )
        module_output.weight =
            torch.nn.Parameter(module.weight.unsqueeze(-1).repeat(1,1,1,1,module.kernel_size[0]))
```

3D segmentation (note read the comment to adjust the code) :

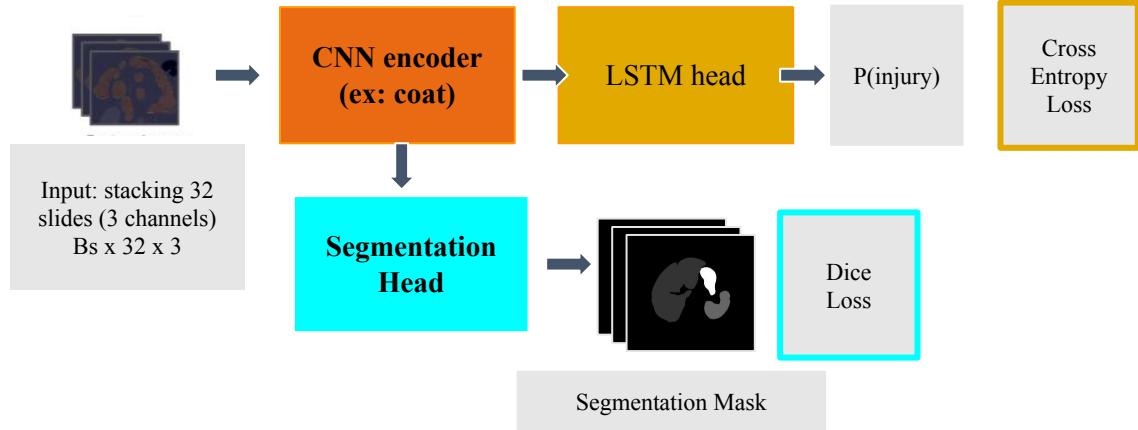
<https://www.kaggle.com/code/haqishen/rsna-2023-1st-place-solution-train-3d-seg/notebook>

2.5D Model

□ Feature extraction and LSTM head



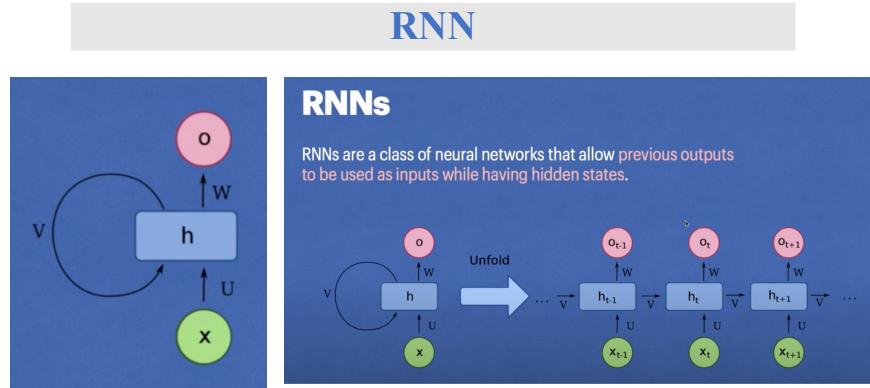
**Stage 2 - 2.5D CNN + RNN
Using 2D segmentation (auxiliary)**



- **CNN model** can be replaced by ViT model (Swin, CoaT...)
- **RNN** can be replaced RNN series (RNN, LSTM) + Attention / Transformer Head

2.5D Model

□ Review RNN/LSTM



Source: [PyTorch RNN Tutorial - Name Classification Using A Recurrent Neural Net](#)

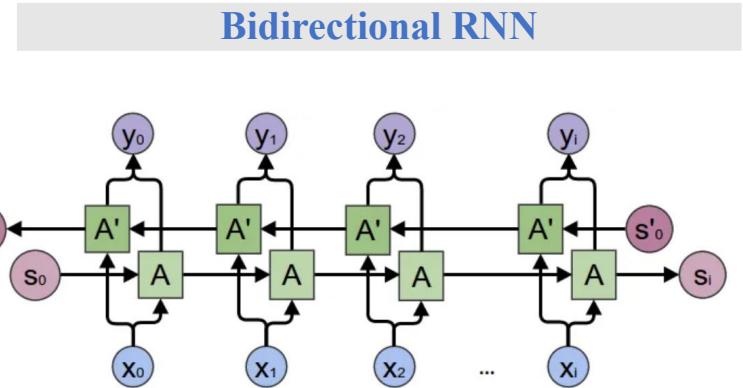
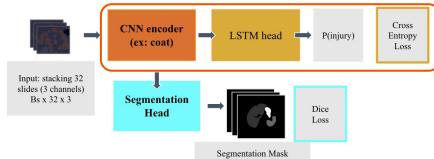


Fig 1: General Structure of Bidirectional Recurrent Neural Networks. Source: colah's blog

2.5D CNN + Sequences Model

□ Feature extraction and LSTM head | Simple Step by Step version



```
import timm
import torch
import torch.nn as nn

# Generate a dummy input tensor
input_tensor = torch.randn(2, 11, 3, 224, 224)
print("- Original input size:", input_tensor.size())

# Define model name and create encoder
name = 'coatnet_1_rw_224'
encoder = timm.create_model(
    name,
    pretrained=None,
    num_classes=0,
    global_pool="avg" if "coat" in name else ""
)

# Reshape input tensor for CNN
bs, n_frames, c, h, w = input_tensor.size()
x = input_tensor.view(bs * n_frames, c, h, w)
print("- Reshaped input for CNN:", x.shape,
      f"\n batch size = {bs * n_frames}",
      f"num channels = {c}, height and width is {h}x{w}")

# Pass the reshaped input through the encoder
features = encoder(x)
num_classes = 11

print("- Output after encoder:", features.shape)
```

- Original input size: torch.Size([2, 11, 3, 224, 224])
- Reshaped input for CNN: torch.Size([22, 3, 224, 224])
batch size = 22 num channels = 3, height and width is 224x224
- Output after encoder: torch.Size([22, 768])

```
# Define LSTM layer

lstm_layer = nn.LSTM(
    input_size=encoder.num_features,
    hidden_size=encoder.num_features // 4,
    # we will take the output mean and max → combine
    batch_first=True,
    bidirectional=True
)
# Reshape features for LSTM input
features = features.view(bs, n_frames, -1)
print("- Reshaped features for LSTM:", features.shape)
print("- After compute mean: ", features.shape)
lstm_output, (hn, cn) = lstm_layer(features)
print("\n LSTM result shape")
print("+ Output size after LSTM layer:", lstm_output.shape)
```

| Reshaped features for LSTM: torch.Size([2, 11, 768])
- After compute mean: torch.Size([2, 11, 768])

```
LSTM result shape
+ Output size after LSTM layer: torch.Size([2, 11, 384])
```

```
mean_ = lstm_output.mean(1)
max_ = lstm_outputamax(1)
print("- Mean:", mean_.shape)
print("- Max:", max_.shape)
features_after_lstm = torch.cat([mean_, max_], -1)
print("- After LSTM mean max:", features_after_lstm.shape)
```

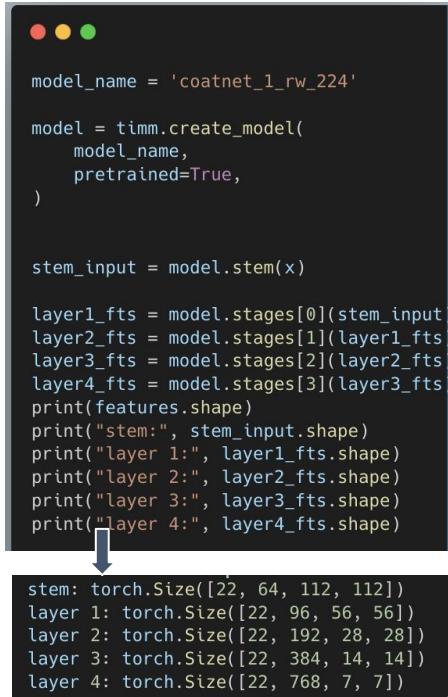
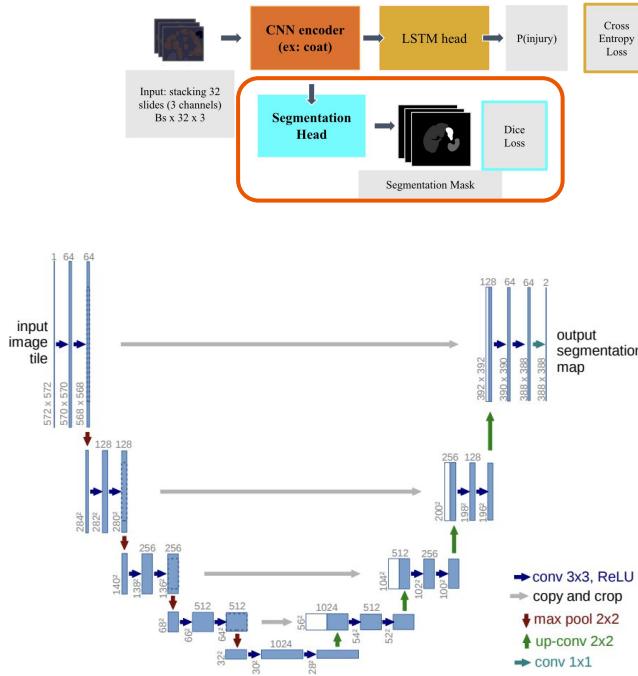
```
linear_layer = nn.Linear(encoder.num_features, 13)
cls_result = linear_layer(features_after_lstm)
print(" \n \n >> Final output logit", cls_result.shape)
```

- Mean: torch.Size([2, 384])
- Max: torch.Size([2, 384])
- After LSTM mean max: torch.Size([2, 768])

```
>> Final output logit torch.Size([2, 13]) |
```

2.5D CNN + Sequences Model

Segmentation Head | Simple Step by Step version



```
model_name = 'coatnet_1_rw_224'

model = timm.create_model(
    model_name,
    pretrained=True,
)

stem_input = model.stem(x)

layer1_fts = model.stages[0](stem_input)
layer2_fts = model.stages[1](layer1_fts)
layer3_fts = model.stages[2](layer2_fts)
layer4_fts = model.stages[3](layer3_fts)
print(features.shape)
print("stem:", stem_input.shape)
print("layer 1:", layer1_fts.shape)
print("layer 2:", layer2_fts.shape)
print("layer 3:", layer3_fts.shape)
print("layer 4:", layer4_fts.shape)

encoder_channels = [64, 96, 192, 384, 768]
decoder_channels = [384, 192, 96, 64]

unet_block_4 = UNetDecoderBlock(
    in_channels = 768, res_channels = 384, out_channels = 384)
unet_fts_4 = unet_block_4(layer4_fts, layer3_fts)
print("unet_fts_4:", unet_fts_4.shape)

unet_block_3 = UNetDecoderBlock(
    in_channels = 384, res_channels = 192, out_channels = 192)
unet_fts_3 = unet_block_3(layer3_fts, layer2_fts)
print("unet_fts_3:", unet_fts_3.shape)

unet_block_2 = UNetDecoderBlock(
    in_channels = 192, res_channels = 96 , out_channels = 96)
unet_fts_2 = unet_block_2(layer2_fts, layer1_fts)
print("unet_fts_2:", unet_fts_2.shape)

unet_block_1 = UNetDecoderBlock(
    in_channels = 96, res_channels = 64 , out_channels = 64)
unet_fts_1 = unet_block_1(layer1_fts, stem_input)
print("unet_fts_1:", unet_fts_1.shape)

features_multiple_resolutions = [
    unet_fts_4, unet_fts_3, unet_fts_2, unet_fts_1]
```

2.5D CNN + Sequences Model

□ Segmentation Head - use FPN to combine features

```
import torch
import torch.nn as nn
import torch.nn.functional as F

class FPN(nn.Module):
    def __init__(self, input_channels: list, output_channels: list, mask_size = (112, 112)):
        super(FPN, self).__init__()
        self.mask_size = mask_size
        self.convs = nn.ModuleList([
            nn.Sequential(
                nn.Conv2d(in_ch, out_ch * 2, kernel_size=3, padding=1), # 384, 64 (batch size =
2*11=22frames)
                nn.ReLU(inplace=True),
                nn.BatchNorm2d(out_ch * 2), # Changed to BatchNorm2d
                nn.Conv2d(out_ch * 2, out_ch, kernel_size=3, padding=1) # 64, 32
            ) for in_ch, out_ch in zip(input_channels, output_channels)]
        )

    def forward(self, xs: list):
        # Scale feature maps to the same resolution
        hcs = [
            F.interpolate(c(x), size=self.mask_size, mode='bilinear', align_corners=False)
            for i, (c, x) in enumerate(zip(self.convs, xs))
        ]
        print(len(hcs))
        return torch.cat(hcs, dim=1) # Concatenate along channel dimension
```

□ Segmentation Head - all together!

```
class SegmentationHead(nn.Module):
    def __init__(self, encoder, mask_size=(112, 112)):
        super(SegmentationHead, self).__init__()
        self.encoder = encoder
        self.mask_size = mask_size

    if 'coat' in self.encoder.name:
        encoder_channels = [64, 96, 192, 384, 768]
        decoder_channels = [384, 192, 96, 64]
        output_channels = 32
        unet_blocks = nn.ModuleList([
            UNetDecoderBlock(encoder_channels[-(i + 1)], encoder_channels[-(i + 2)],
decoder_channels[i])
            for i in range(len(decoder_channels))
        ])
        self.fpn = FPN(decoder_channels, [output_channels * 4])

        current_channels = output_channels * 3
        num_layers = len(decoder_channels) # Adjust as needed
        intermediate_channels = [current_channels // 2**l for l in range(num_layers)]
        layers = []
        for ic in intermediate_channels:
            layers.append(nn.Conv2d(current_channels, ic, kernel_size=3, stride=2, padding=1))
            layers.append(nn.BatchNorm2d(ic))
            layers.append(nn.ReLU(inplace=True))
            current_channels = ic
        layers.append(nn.Conv2d(current_channels, 1, kernel_size=1))
        self.seg_cnn_layers = nn.Sequential(*layers)

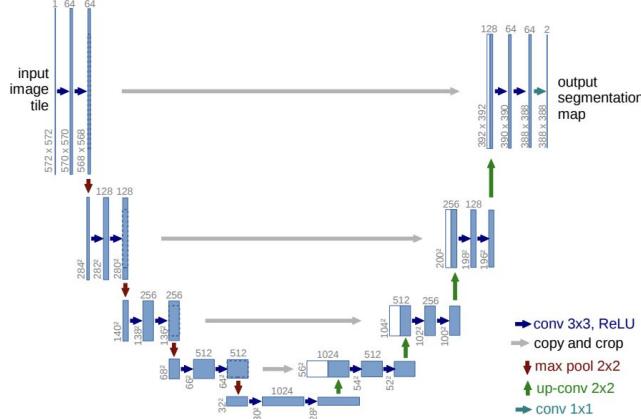
    def forward(self, intermediate_features):
        if 'coat' in self.encoder.name:
            layer4_encoder_feature, layer3_encoder_feature, layer2_encoder_feature,
layer1_encoder_feature, stem_input = intermediate_features

            unet_fts_4 = self.unet_blocks[3](layer4_encoder_feature, layer3_encoder_feature)
            unet_fts_3 = self.unet_blocks[2](layer3_encoder_feature, layer2_encoder_feature)
            unet_fts_2 = self.unet_blocks[1](layer2_encoder_feature, layer1_encoder_feature)
            unet_fts_1 = self.unet_blocks[0](layer1_encoder_feature, stem_input)

            segmentation_features = self.fpn([unet_fts_4, unet_fts_3, unet_fts_2, unet_fts_1])
            x = F.interpolate(
                segmentation_features, size=self.mask_size, mode='bilinear', align_corners=False)
            pred_masks = self.seg_cnn_layers(x)
            return pred_masks
        else:
            return None |
```

2.5D CNN + Sequences Model

Segmentation Head - all together!



```
class SegmentationHead(nn.Module):
    def __init__(self, encoder, mask_size=(112, 112)):
        super(SegmentationHead, self).__init__()
        self.encoder = encoder
        self.mask_size = mask_size

    if 'coat' in self.encoder.name:
        encoder_channels = [64, 96, 192, 384, 768]
        decoder_channels = [384, 192, 96, 64]
        output_channels = 32
        self.unet_blocks = nn.ModuleList([
            UNetDecoderBlock(encoder_channels[-(i+1)], encoder_channels[-(i+2)],
            decoder_channels[i])
            for i in range(len(decoder_channels)))
        ])
        self.fpn = FPN(decoder_channels, [output_channels * 4])

    current_channels = output_channels * 3
    num_layers = len(decoder_channels) # Adjust as needed
    intermediate_channels = [current_channels // 2**i for i in range(num_layers)]
    layers = []
    for ic in intermediate_channels:
        layers.append(nn.Conv2d(current_channels, ic, kernel_size=3, stride=2, padding=1))
        layers.append(nn.BatchNorm2d(ic))
        layers.append(nn.ReLU(inplace=True))
        current_channels = ic
    layers.append(nn.Conv2d(current_channels, 1, kernel_size=1))
    self.seg_cnn_layers = nn.Sequential(*layers)

    def forward(self, intermediate_features):
        if 'coat' in self.encoder.name:
            layer4_encoder_feature, layer3_encoder_feature, layer2_encoder_feature,
            layer1_encoder_feature, stem_input = intermediate_features

            unet_fts_4 = self.unet_blocks[3](layer4_encoder_feature, layer3_encoder_feature)
            unet_fts_3 = self.unet_blocks[2](layer3_encoder_feature, layer2_encoder_feature)
            unet_fts_2 = self.unet_blocks[1](layer2_encoder_feature, layer1_encoder_feature)
            unet_fts_1 = self.unet_blocks[0](layer1_encoder_feature, stem_input)

            segmentation_features = self.fpn([unet_fts_4, unet_fts_3, unet_fts_2, unet_fts_1])
            x = F.interpolate(
                segmentation_features, size=self.mask_size, mode='bilinear', align_corners=False)
            pred_masks = self.seg_cnn_layers(x)
            return pred_masks
        else:
            return None
```

2.5D CNN + Sequences Model

□ 2.5D CNN + Sequence Model - all together!

```
class _25DCnnRnnSegAux(nn.Module):
    ...
    2.5 D model : CNN model + RNN with Segmentation head
    ...
    def __init__(self,
                 encoder,
                 num_classes = 2,
                 num_classes_aux = 11,
                 n_channels = 3,
                 head_3d = "",
                 n_frames = 1,
                 dropout_rate=0,
                 segmentation_aux=True):
        super().__init__()
        #set up
        self.encoder = encoder
        self.num_features=encoder.num_features
        self.num_classes = num_classes
        self.num_classes_aux = num_classes_aux
        self.n_channels = n_channels
        self.head_3d = head_3d
        self.n_frames = n_frames
        self.dropout = nn.Dropout(p=dropout_rate) if dropout_rate > 0 else nn.Identity()
        # layers.
        if head_3d == 'lstm':
            self.lstm = nn.LSTM(input_size = self.num_features,
                                hidden_size = self.num_features // 4,
                                batch_first = True,
                                bidirectional=True)

        self.last_layer = nn.Linear(self.num_features, self.num_classes)

        if self.num_classes_aux > 0:
            if segmentation_aux:
                self.aux_layer = SegmentationHead(self.encoder)
            else:
                self.aux_layer = nn.Linear(self.num_features, self.num_classes_aux)
```

```
def extract_features(self, x):
    feature_output = self.encoder(x)
    intermediate_feature = None
    if self.segmentation_aux and 'coat' in self.encoder.name:
        layer0_encoder_feature = self.encoder.stems[0]
        layer1_encoder_feature = self.encoder.stages[0](layer0_encoder_feature)
        layer2_encoder_feature = self.encoder.stages[1](layer1_encoder_feature)
        layer3_encoder_feature = self.encoder.stages[2](layer2_encoder_feature)
        layer4_encoder_feature = self.encoder.stages[3](layer3_encoder_feature)

    intermediate_features = [
        layer4_encoder_feature,
        layer3_encoder_feature,
        layer2_encoder_feature,
        layer1_encoder_feature,
        layer0_encoder_feature
    ]

    fts = self.dropout(fts)
    return fts, intermediate_features

def forward_head3d(self, x):
    if self.head_3d == 'lstm':
        x, _ = self.lstm(x)
        mean = x.mean(1)
        max_ = x.amax(1)
        x = torch.cat([mean, max_], dim=1)
    return x
```

```
def forward(self, x):
    if self.head_3d:
        bs, n_frames, c, h , w = x.size()
        x = x.view(bs*n_frames, c, h, w)

    fts, intermediate_features = self.extract_features(x)

    print("fts size: ", fts.shape)

    if self.head_3d != "":
        fts = fts.view(bs, n_frames, -1)
        fts = self.forward_head3d(fts)

    output = self.last_layer(fts)

    if self.num_classes_aux:
        segmentation_output = self.segmentation_head(intermediate_features)
    else:
        segmentation_output = None
    if self.num_classes_aux > 0:
        aux_output = self.aux_layer(fts)
    else:
        aux_output = torch.zeros((fts.size(0),))

    # processing column mask into slice -> input to the model

    return output, aux_output |
```

```
1 class CustomLoss(nn.Module):
2     def __init__(self):
3         super(CustomLoss, self).__init__()
4         #self.bce = nn.BCEWithLogitsLoss(pos_weight=torch.as_tensor([2.318]).cuda())
5         self.bce = nn.BCEWithLogitsLoss()
6         self.dice = smp.losses.DiceLoss(smp.losses.MULTILABEL_MODE, from_logits=True)
7
8     def forward(self, outputs, targets, masks_outputs, masks_targets):
9         loss1 = self.bce(outputs, targets)
10
11         masks_outputs = masks_outputs.float()
12         # masks_outputs2 = masks_outputs2.float()
13
14         masks_targets = masks_targets.float().flatten(0, 1)
15
16         loss2 = self.dice(masks_outputs, masks_targets) #+ self.dice(masks_outputs2,
17
18
19         loss = loss1 + (loss2 * CFG.segw)
20
21     return loss
```

Upsampling layer (optional)

❑ Upsampling layer: ConvTranspose2d, Interpolation or Pixel Shuffle

ConvTranspose2d

Input Kernel

$$\begin{array}{|c|c|} \hline 0 & 1 \\ \hline 2 & 3 \\ \hline \end{array} \quad \text{Transposed Conv} \quad \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 2 & 3 \\ \hline \end{array}$$

Output

$$= \begin{array}{|c|c|} \hline 0 & 0 \\ \hline 0 & 0 \\ \hline \end{array} + \begin{array}{|c|c|} \hline 0 & 1 \\ \hline 2 & 3 \\ \hline \end{array} + \begin{array}{|c|c|} \hline 0 & 2 \\ \hline 4 & 6 \\ \hline \end{array} + \begin{array}{|c|c|} \hline 0 & 3 \\ \hline 6 & 9 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 0 & 0 & 1 \\ \hline 0 & 4 & 6 \\ \hline 4 & 12 & 9 \\ \hline \end{array}$$

Figure: ConvTranspose2d 2x2
Src: D2l.ai



Using deconvolution.
Heavy checkerboard artifacts.

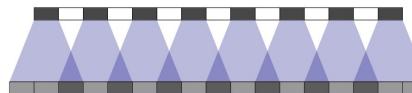
Artifact



Using resize-convolution.
No checkerboard artifacts.

Less Artifact

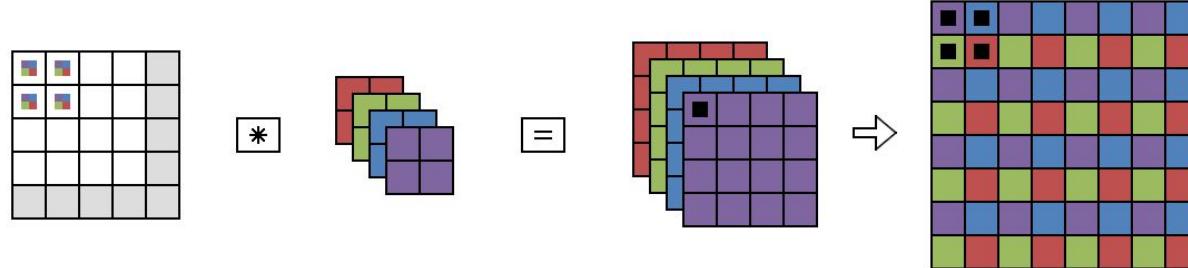
The overlap pattern also forms in two dimensions. The uneven overlaps on the two axes multiply together, creating a characteristic checkerboard-like pattern of varying magnitudes.



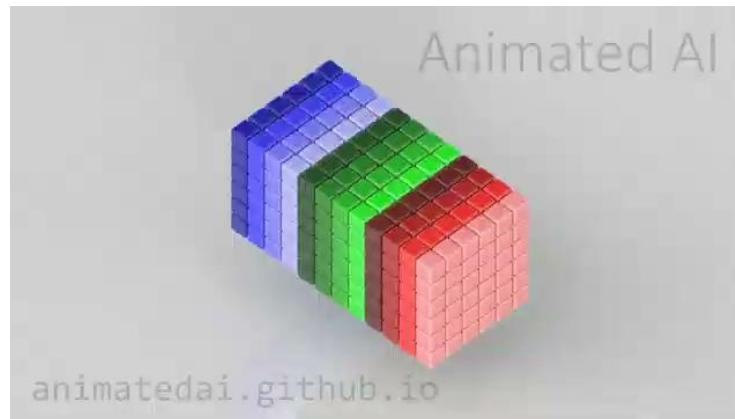
Source: <https://distill.pub/2016/deconv-checkerboard/>

Upsampling layer (optional)

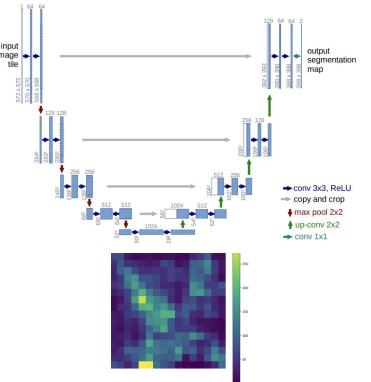
Pixel Shuffle



Source: <https://stackoverflow.com/questions/72069125/is-convolution-pixelshuffle-the-same-as-subpixel-convolution>



Upsampling layer (optional)



Upsampled Image (ConvTranspose2d)



Upsampled Image (Interpolate)



Upsampled Image (Pixel Shuffle)



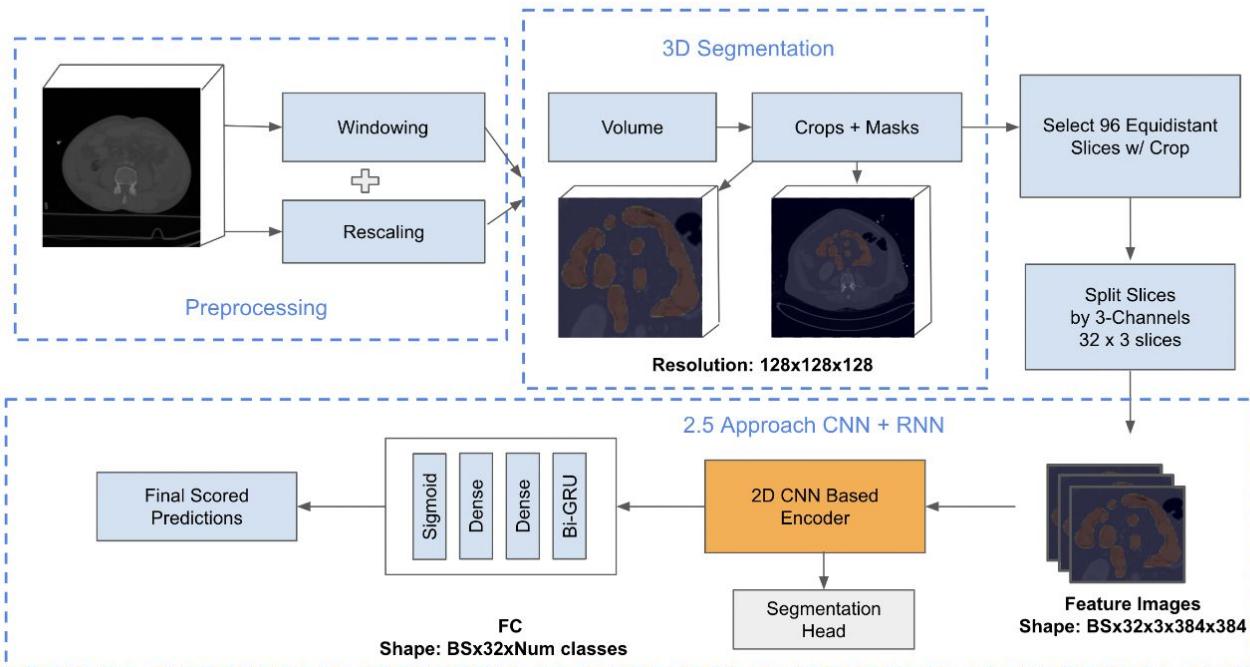
Upsampled Image
(PixelShuffle with ICNR initialization)



Notebook Link: [\[vizualize\] pixel-shuffle vs conv-transpose](#)

Review pipeline

2 Stage Architecture



Summary

Volumetric data

3D CNN

3D Unet

Exploring RSNA with 3D Unet and 2.5D CNN + RNN



RSNA RADIOLOGICAL SOCIETY OF NORTH AMERICA - FEATURED CODE COMPETITION - 2 MONTHS TO GO

Submit Prediction ...

RSNA 2024 Lumbar Spine Degenerative Classification

Classify lumbar spine degenerative conditions

Thank you for watching !

Let's learn more together!



AI VIET NAM
[@aivietnam.edu.vn](http://aivietnam.edu.vn)

